

Operating Systems Structure

- How do user programs obtain services from the OS in a safe way?
- By what principles are OSes designed?

COSC 301: Operating Systems
10 September 2012

Joel Sommers
jsommers@colgate.edu
Colgate University

Operating
Systems
Structure

What is an
operating system?
Hardware
System Calls
Designing an OS

10 September 2012 (1/18)

Learning Goals

- ▶ Define the roles and purposes of an operating system.
- ▶ Describe and contrast user mode and kernel mode in an operating system.
- ▶ Identify characteristics of how and describe why operating systems are structured as monolithic, microkernel, and virtualized systems.
- ▶ Effectively distinguish between “mechanism” and “policy” with respect to the design of an operating system.
- ▶ Describe what a system call is, and sketch the steps that take place to carry out a system call.

Operating
Systems
Structure

What is an
operating system?
Hardware
System Calls
Designing an OS

10 September 2012 (2/18)

The OS as a Virtual Machine

- ▶ Physical machine
 - ▶ Runs 1 program on the CPU at a time
 - ▶ Has 1 physical memory
 - ▶ Has devices (like disks) that are really hard to use
- ▶ OS provides virtual machine illusion to programs
 - ▶ Many CPUs, separate and large memories for each program, easy use of devices
- ▶ Key issues: efficiency, protection

Operating
Systems
Structure

What is an
operating system?
Hardware
System Calls
Designing an OS

10 September 2012 (3/18)

What is the role of the OS?

Role #1: Provide standard library to applications

- ▶ What is a *resource*?
 - ▶ Anything computationally valuable (e.g., CPU, memory, disk, network)
- ▶ Advantages of a standard library
 - ▶ Allow applications to reuse common facilities
 - ▶ Make different devices appear the same
 - ▶ Provide higher-level abstractions
- ▶ Challenges
 - ▶ What are the right abstractions?
 - ▶ How much of the hardware should be exposed?

Operating
Systems
Structure

What is an
operating system?
Hardware
System Calls
Designing an OS

10 September 2012 (4/18)

What is the role of the OS?

Operating
Systems
Structure

What is an
operating system?

Hardware

System Calls

Designing an OS

Role #2: Resource coordinator and allocator

- ▶ Advantages of a resource coordinator
 - ▶ Virtualize resources so multiple users or applications can share
 - ▶ Protect applications from one another
 - ▶ Provide efficient and fair access to resources
 - ▶ Can arbitrate among conflicting requests
- ▶ Challenges
 - ▶ What are the right *mechanisms*?
 - ▶ What are the right *policies*?

10 September 2012 (5/18)

What functionality belongs in the OS?

Operating
Systems
Structure

What is an
operating system?

Hardware

System Calls

Designing an OS

No single right answer

- ▶ Desired functionality depends on outside factors
- ▶ OS must adapt to both user expectations and technology changes
 - ▶ Change abstractions provided to users
 - ▶ Change algorithms used to implement those abstractions
 - ▶ Change low-level implementation to deal with hardware
- ▶ Current operating systems driven by technology and application trends
- ▶ OS design is (almost always) about tradeoffs (performance, usability, security, generality, . . .)

10 September 2012 (6/18)

OS components

Operating
Systems
Structure

What is an
operating system?

Hardware

System Calls

Designing an OS

Definition

Kernel: core components of the OS

- ▶ Process scheduler
 - ▶ **process: a running program**
 - ▶ Determines when and for how long each process executes
- ▶ Memory manager
 - ▶ Determines when and how memory is allocated to processes
 - ▶ Decides what to do when main memory is full
- ▶ File system
 - ▶ Organizes named collections of data in persistent storage
- ▶ Networking
 - ▶ Enables processes to communicate with one another

10 September 2012 (7/18)

User Program and OS Interaction

Operating
Systems
Structure

What is an
operating system?

Hardware

System Calls

Designing an OS

- ▶ The OS kernel manages hardware resources, provides services via a standard set of libraries to user programs
 - ▶ User applications make **system calls** via the OS libraries to request kernel services
 - ▶ Certain CPU hardware features enable the OS to do this in a safe way
 - ▶ Example: user programs shouldn't be allowed to use the `halt` instruction without proper permissions
 - ▶ Example: user programs shouldn't be allowed to directly access hardware resources unless explicitly allowed by the OS
 - ▶ How the OS interacts with hardware devices depends on the goals of the OS, and capabilities of the hardware
 - ▶ The specific communication and/or software mechanisms used by the user applications to obtain library services depend on **OS kernel design**

10 September 2012 (8/18)

Some important CPU features

Standard stuff:

- ▶ Basic cycle: fetch, decode, execute
- ▶ Instructions: load, store, register ops
- ▶ Superscalar pipelining

Special-purpose registers:

- ▶ Program counter (instruction pointer)
- ▶ Stack pointer (top of current stack)
- ▶ PSW (program status word)
 - ▶ Condition code bits (results of comparison instructions)
 - ▶ **Mode bit(s)**: enables the OS to control access to hardware
 - ▶ User mode: only a subset of instructions/features are accessible
 - ▶ Kernel mode: all instructions and features are accessible
 - ▶ Some CPUs support more than 2 levels of protection

Operating
Systems
Structure

What is an
operating system?

Hardware

System Calls

Designing an OS

10 September 2012 (9/18)

Context switching

- ▶ Whenever execution switches between protection domains, a **context switch** occurs
- ▶ Actions on context switch:
 - ▶ Save the PC, stack pointer, PSW (as well as any other vital CPU state)
 - ▶ Save the contents of all general purpose registers
 - ▶ Save and reprogram the memory-management unit registers (more later on what this is all about)
 - ▶ Wait while the instructions currently in the CPU pipeline are trashed
 - ▶ (Wait for instruction/data caches to load from new program's memory; details depend heavily on CPU architecture)
- ▶ Context switches can be expensive. . .

Operating
Systems
Structure

What is an
operating system?

Hardware

System Calls

Designing an OS

10 September 2012 (10/18)

Polling and Interrupts

- ▶ Two fundamental ways for an OS to interact with hardware and perform I/O: **interrupts** and **polling**

Interrupts:

- ▶ Hardware device interrupts CPU from its current task
 - ▶ Example: key is pressed on the keyboard
 - ▶ Example: packet arrives at a network interface
- ▶ OS maintains a list of routines to handle hardware I/O events (**interrupt vector**)
 - ▶ Upon hardware interrupt, OS saves CPU state, switches to kernel mode, uses device id that generated to interrupt to find the interrupt routine to execute
 - ▶ Interrupt routine calls into device driver to perform I/O
 - ▶ Control is eventually returned to previously executing process at the next instruction
- ▶ Interrupt-driven I/O is very commonly used in commodity hardware

Operating
Systems
Structure

What is an
operating system?

Hardware

System Calls

Designing an OS

10 September 2012 (11/18)

Polling

- ▶ OS checks whether there is I/O to be done on a particular device
- ▶ May check repeatedly in a tight loop if it expects input, or may check periodically or infrequently
 - ▶ Depends on OS goals, application requirements
- ▶ **Polling is a common scenario in real-time systems that cannot suffer unpredictable delays due to interrupts**
 - ▶ It can also improve performance when heavy I/O traffic is expected, e.g., a web server that gets a lot of network traffic
- ▶ Hardware must support polling; not all commodity hardware does (need capability to disable interrupts)

Operating
Systems
Structure

What is an
operating system?

Hardware

System Calls

Designing an OS

10 September 2012 (12/18)

How does a user process ask the OS to perform a function for it?

- ▶ Scenario:
 - ▶ User program is running with CPU mode in user mode
 - ▶ User program needs to obtain information from system hardware, or change the state of system hardware (e.g., perform I/O), or obtain/modify information maintained by the OS (e.g., information about other system processes)
- ▶ Problem: user program cannot directly access kernel memory or hardware with the CPU running in user-mode
- ▶ Must make a **system call** to ask the OS to perform some action on its behalf
 - ▶ e.g., `fopen()`, `getrusage()`, `fgets()`, ...
 - ▶ Can use `strace` to trace syscalls in any executable

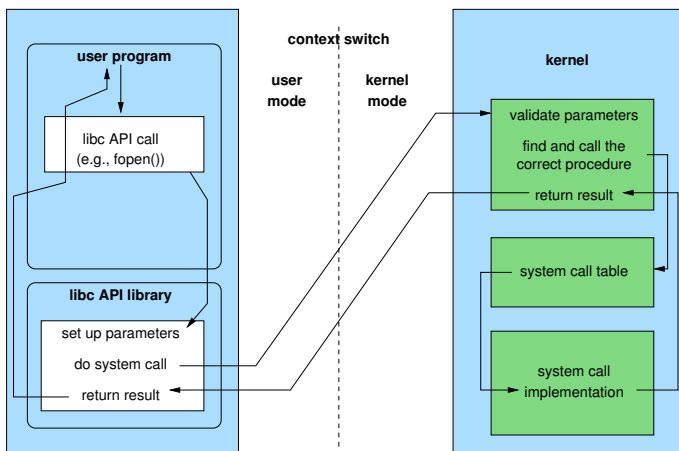
10 September 2012 (13/18)

Traps

- ▶ A system call is the mechanism by which a user process calls a kernel procedure
 - ▶ Used to do I/O, read/write files, etc.
- ▶ Historically, used a similar mechanism as interrupts
 - ▶ User program generates a trap
 - ▶ Kernel exception handler is called (similar to interrupt handler)
 - ▶ OS identifies which system call is required, and calls the relevant procedure
 - ▶ Returns execution to user process on completion
- ▶ INT instruction in x86 used to invoke software traps
- ▶ 8-bit opcode allows 256 different exceptions
 - ▶ Linux: INT 0x80 is a system call
 - ▶ Windows: INT 0x2E is a system call
- ▶ Modern mechanism: `syscall/sysenter` instruction (eliminates much of the overhead of a traditional interrupt)

10 September 2012 (14/18)

Processing a system call



Typical system call processing sequence

10 September 2012 (15/18)

OS kernel design

- ▶ How should kernel subsystems communicate with one another?
- ▶ Should all kernel subsystems execute in kernel mode or user mode?

Monolithic kernel Communication through function calls; everything runs in the same privileged protection ring

- ▶ Pro: fast
- ▶ Con: one part of the kernel crashes, and the whole system goes down

Microkernel Small core kernel component runs in privileged mode; kernel subsystems run in user mode; communication among components via message passing

- ▶ Pro: typically only core kernel crash can take the system down (isolation)
- ▶ Con: overhead of message passing

10 September 2012 (16/18)

Modularization

- ▶ Dynamically loadable modules implement various OS subsystems
 - ▶ Device drivers are the most common
 - ▶ Can also be more fundamental components of the OS, *e.g.*, CPU scheduler
- ▶ Pro: Running OS only contains desired/required features; functionality can easily be modified by swapping modules that implement a particular feature
- ▶ Con: A little bit harder to implement (far outweighed by pros)
- ▶ Somewhat orthogonal issue to kernel communication and isolation design
 - ▶ `lsmod` on Linux to see what kernel modules are loaded

A key design goal for an OS is to separate **policy** and **mechanism**.

- ▶ OS kernel should avoid embedding policies
- ▶ Provide neutral low-level mechanisms upon which higher layers can implement various policies
- ▶ Separation allows policies to change, perhaps on-the-fly

Policy or mechanism?

- ▶ Saving the register state of a process
- ▶ How long a user process is allowed to execute on the CPU
- ▶ The timer interrupt
- ▶ Continuing to run the current process when a disk I/O interrupt occurs