

# COSC 301: Weds-Fri, 12-14 Sept 2012

## Learning goals

- Identify the attributes and purposes of a process.
- Describe the actions that take place on a context switch.

## Virtualizing the CPU

Goal: run  $N$  processes “at once” even through only  $M$  CPUs ( $N \gg M$ )

- This is termed “multiprogramming”
  - Multiple programs residing in memory at once, possibly running simultaneously
- Each process thinks it has its own isolated machine
- Also, each thinks it has its own *restricted* and *efficient* machine
  - That is, the cost of creating the illusion shouldn’t be high
  - And we want to be sure to protect/isolate processes from each other
- **Why?** Why do we want to do any of this?

---

Key abstraction: the process (a running program)

- What does a running program do?
  - Has a private memory (“address space”)
  - Simple operations on memory and registers
    - \* Virtualizing the CPU will require us to (in a sense) virtualize the registers!
- To support the process abstraction, need
  - Mechanisms** *How* to virtualize Name of general mechanism: “limited direct execution”
  - Policies** Which process to run and how long? (scheduling, next week)

---

How to virtualize CPU?

- Start simple: when you want to run a program, just run it
  - Load it into memory, point PC at first instruction, go!
  - This is *direct execution*
- What’s wrong? What are the problems that can arise?
  1. What if process wants to do something that is restricted?
  2. What if OS wants to stop process A and run process B?
  3. What if the running process does something *slow*, like I/O?

---

**Problem 1:** Restricted operations (like I/O, device operations, etc.)

We could trust processes not to read the files of other processes, but that's probably just a little bit naive.

- Solution: operating system privilege levels
  - Process can be running in one of two modes:
    1. “User” mode (restricted)
    2. “Kernel” mode (not restricted)
  - Need a way to change between these two modes: the **system call**
  - Scenario: process runs, calls syscall (e.g., `fopen`), causes trap (implicit change in privilege), does restricted op, returns.
- On `trap` instruction (`INT` on x86)
  1. Save state: done by hardware (sometimes also software)
    - logically, save PC, registers to a special place, then run OS code w/o worry of corrupting user registers
    - why must the saving of state be done (at least partially) by hardware?
    - OS issue: where to put the saved state? (don't answer Q yet...)
  2. Raise privilege level - allows us to do some fun stuff
  3. Jump to correct kernel code (trap handler, or interrupt handler)
    - How does hardware know which PC to jump to?
    - Protocol:
      - \* Kernel boots (into kernel mode)
      - \* Installs trap tables (list of addresses of functions to run when certain traps/interrupts happen)
      - \* OS eventually runs a process (in user mode)
    - Process makes sys calls
    - So on trap, the hardware consults the trap table (interrupt vector) to know which code to run
- On return from trap
  1. Lower privilege level
  2. Restore state of registers of calling process
  3. Jump back to PC after trap in user code

What if a process running in user mode tries to do something restricted? E.g., call `halt` instruction?

- Typically the OS will kill it

---

**Problem 2:** How to stop a currently running process A and run process B? (Assume only 1 CPU: when A is running, B cannot be running)

**Cooperative** just wait for process to call into kernel or yield explicitly

**Non-cooperative**

forcibly reclaim the CPU via a **timer interrupt**

- turn on timer interrupt at boot, before *any* user process
- goes off every x milliseconds
- at timer interrupt, decide to continue A or switch to B

What's the main issue for deciding whether we should use *cooperative* or *non-cooperative* scheduling? (trust!)

If we decide to switch the running process: **context switch**

1. Save A state, stop it

2. Restore B's state, run it

OS issues

**trust** processes may misbehave or may be buggy

**bookkeeping** need to keep track of state of each process (whether running, or ready to be run)

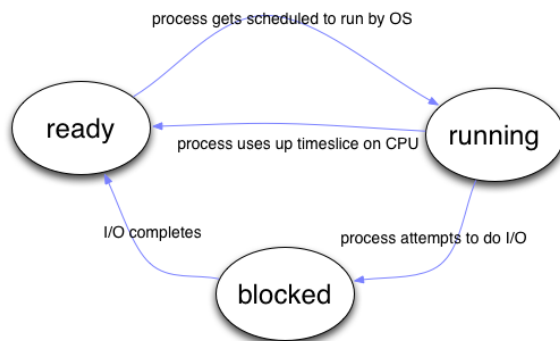
need a data structure to track this info: *process list*

- Per-process info stored: saved registers (answers question from above on where we save process state on a syscall), state: ready/running

---

### Problem 3: Slow operations like I/O

- Basic solution: keep more detailed information about process state



---

Summary: Limited Direct Execution

**direct** Want process to run on the CPU directly as much as possible

**limited** At certain key points, OS and hardware impose themselves to regain/retain control of the system

Result: a virtualized CPU

- With safety (why?)
- With efficiency (why?)