

Virtualizing Memory

- Memory virtualization goals
- Address translation and dynamic relocation
- Segmentation
- Dealing with fragmentation

COSC 301: Operating Systems
28 September 2012

Joel Sommers
jsommers@colgate.edu
Colgate University

Virtualizing
Memory

Memory
Virtualization
Dynamic
relocation
Segmentation
Fragmentation

28 September 2012 (1/21)

Learning goals

- ▶ Describe a base and bounds-based dynamic memory relocation scheme for supporting multiprogramming.
- ▶ Compare simple dynamic memory relocation and segmentation as approaches for allocating memory to user processes.
- ▶ Describe and compare algorithms for memory allocation, e.g., first, next, best fit, and the buddy algorithm.

Virtualizing
Memory

Memory
Virtualization
Dynamic
relocation
Segmentation
Fragmentation

28 September 2012 (2/21)

Memory virtualization goals

Main goal: extend limited direct execution to memory references to create illusion of a private address space for each process

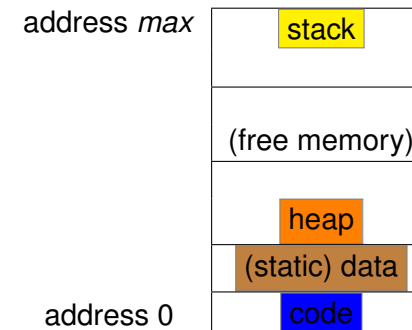
- ▶ Mechanism: address translation
- ▶ Want it to be **efficient**, **transparent**, and **controlled** (protected)
- ▶ What about **isolation**?

Virtualizing
Memory

Memory
Virtualization
Dynamic
relocation
Segmentation
Fragmentation

28 September 2012 (3/21)

Process Logical Address Space




- ▶ **logical** address space == **virtual** address space
- ▶ Linear address space from address 0 to *max*
- ▶ **Purely conceptual view of a process's memory**
 - ▶ Physical memory address space \neq logical process address space

Virtualizing
Memory

Memory
Virtualization
Dynamic
relocation
Segmentation
Fragmentation

28 September 2012 (4/21)

Dynamic Relocation



- Assume that P1 base physical address is 1024, base physical address for P2 is 4096.
- Say P1 and P2 each execute `load 100, R1` instruction — what should happen?
 - How to map (translate) *virtual* memory reference to physical address?
 - What if P1 issues instruction `load 4000, R1`?

Virtualizing Memory

Memory Virtualization
Dynamic relocation
Segmentation
Fragmentation

28 September 2012 (5/21)

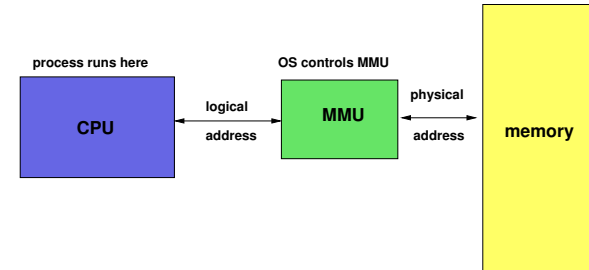
Dynamic Relocation

We need hardware support!

- Memory management unit (MMU)

MMU dynamically changes process address at every memory reference

- Process generates **logical** or **virtual** addresses
- Memory hardware uses **physical** or **real** addresses



Virtualizing Memory

Memory Virtualization
Dynamic relocation
Segmentation
Fragmentation

28 September 2012 (6/21)

Hardware Support for Dynamic Relocation

Two operating modes

- Privileged (protected, kernel) mode: OS
 - When enter OS (trap, system calls, interrupts, exceptions)
 - Allows certain instructions to be executed
 - Can manipulate contents of MMU
 - Allows OS to access all of physical memory
- User mode: user processes
 - Perform translation of logical address to physical address

MMU contains base and bounds registers

- base: start location for process address space
- bounds: size limit of address space

Virtualizing Memory

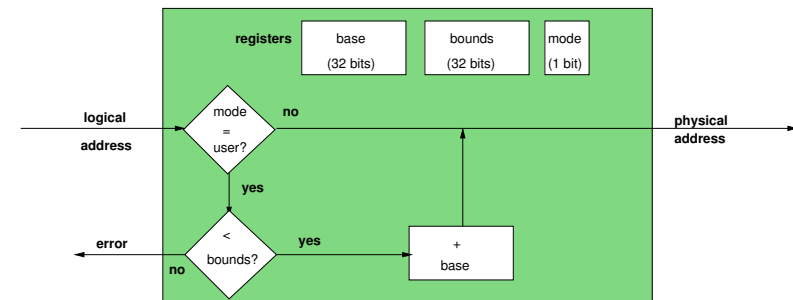
Memory Virtualization
Dynamic relocation
Segmentation
Fragmentation

28 September 2012 (7/21)

Implementation of Dynamic Relocation

Translation on every memory access of user process

- MMU compares logical address to bounds register
 - If logical address is greater, then generate error (kill process)
- MMU adds base register to logical address to form physical address



Virtualizing Memory

Memory Virtualization
Dynamic relocation
Segmentation
Fragmentation

28 September 2012 (8/21)

Example of Dynamic Relocation

What are the physical addresses for the following 16-bit logical addresses?

- ▶ Process 1: base=0x4320, bounds=0x2220
 - ▶ 0x0000:
 - ▶ 0x1110:
 - ▶ 0x3000:
- ▶ Process 2: base=0x8450, bounds=0x3330
 - ▶ 0x0000:
 - ▶ 0x1110:
 - ▶ 0x3000:
- ▶ OS
 - ▶ 0x0000:

Virtualizing
Memory

Memory
Virtualization

Dynamic
relocation

Segmentation

Fragmentation

28 September 2012 (9/21)

OS involvement with Dynamic Relocation

Process creation

- ▶ Initialize base and bounds (how?)

Growth of address space

- ▶ Need to track free/available physical memory
- ▶ Adjust bounds register as necessary, somehow

Context-switch

- ▶ Add base/bounds to PCB
- ▶ Save base/bounds of P1, restore base/bounds of P2

Virtualizing
Memory

Memory
Virtualization

Dynamic
relocation

Segmentation

Fragmentation

28 September 2012 (10/21)

Base and Bounds Discussion

Advantages

- ▶ Fast and simple!
 - ▶ Add and compare can be done in parallel
- ▶ Supports memory virtualization with little overhead
 - ▶ Can change where a process is physically located
 - ▶ **Why might we want to do this?**

Disadvantages

- ▶ Each process must be allocated contiguously in physical memory
 - ▶ Must allocate memory that may not be used by process (how big should the heap be? the stack?)
- ▶ No fine-grain protection beyond protecting complete address spaces
 - ▶ E.g., read/write/execute permissions for different parts of memory
- ▶ No partial sharing; cannot share limited parts of an address space

Virtualizing
Memory

Memory
Virtualization

Dynamic
relocation

Segmentation

Fragmentation

28 September 2012 (11/21)

Segmentation

Basic idea: divide address space into logical segments

- ▶ A *generalization* of dynamic relocation
- ▶ Each segment corresponds to logical entity in address space
 - ▶ Code, stack, heap

Each segment can independently:

- ▶ be placed (and relocated) separately in physical memory
- ▶ grow and shrink
- ▶ be protected (separate read/write/execute protection bits)

Virtualizing
Memory

Memory
Virtualization

Dynamic
relocation

Segmentation

Fragmentation

28 September 2012 (12/21)

Segmented Addressing

Virtualizing
Memory

Memory
Virtualization
Dynamic
relocation

Segmentation
Fragmentation

How does process designate a particular segment?

- ▶ Explicit: use part of logical address
 - ▶ Top bits of logical address select segment
 - ▶ Low bits of logical address select offset within segment

What if small address space, not enough bits?

- ▶ Implicitly by type of memory reference
- ▶ E.g., if instruction fetch, then use code segment

28 September 2012 (13/21)

Segmentation Implementation

Virtualizing
Memory

Memory
Virtualization
Dynamic
relocation

Segmentation
Fragmentation

MMU contains segment table (per process)

- ▶ Each segment has own base and bounds, protection bits
- ▶ Example: 16 bit logical address, max segment address $2^{14} - 1$, 4 segments (2 high bits)

Segment	Base	Bounds	R W X
0	0x2000	0x06ff	1 0 0
1	0x0000	0x04ff	1 0 1
2	0x3000	0x0fff	1 1 0
3	0x0000	0xffff	0 0 0

Translate logical addresses to physical addresses:

- ▶ 0x0240
- ▶ 0x4508
- ▶ 0x865c
- ▶ 0xc002

28 September 2012 (14/21)

Discussion of Segmentation

Virtualizing
Memory

Memory
Virtualization
Dynamic
relocation

Segmentation
Fragmentation

Advantages

- ▶ Enables sparse allocation of address space
 - ▶ Stack and heap can grow independently
- ▶ Different protection for different segments
 - ▶ Read (R), Write (W), Execute (X) for each segment
 - ▶ E.g., RX for code, RW for heap and stack
- ▶ Enables sharing of selected segments
- ▶ Supports dynamic relocation of each segment

Disadvantages

- ▶ Each segment must be allocated contiguously
 - ▶ Can lead to **external fragmentation**
 - ▶ May not have sufficient physical memory for large segments

28 September 2012 (15/21)

Fragmentation

Virtualizing
Memory

Memory
Virtualization
Dynamic
relocation

Segmentation
Fragmentation

Fragment definition: free memory that is too small to be usefully allocated

- ▶ External: visible to allocator
- ▶ Internal: visible to requester (*e.g.*, if must allocate at some granularity; want 8 bytes but can only obtain in 16 byte quantities)

Goal: minimize fragmentation

- ▶ Few holes, each hole is large
- ▶ Free space is contiguous

Fragmentation can occur from the perspective of two different allocators:

- ▶ OS kernel: allocate memory to a process segment
- ▶ Heap allocation: allocate within the heap segment of a process

28 September 2012 (16/21)

Memory allocation algorithms

Best fit

- ▶ Search entire list for each allocation
- ▶ Choose free block that most closely matches size of request
- ▶ Optimization: stop searching if exact match found

First fit

- ▶ Allocate first block that is large enough to fulfill request

Rotating first fit

- ▶ Variant of first fit, remember place in list
- ▶ Start with next free block after most recently allocated block

Worst fit

- ▶ Allocate largest block to request (leaving most left-over space)

Virtualizing
Memory

Memory
Virtualization

Dynamic
relocation

Segmentation

Fragmentation

28 September 2012 (17/21)

Memory allocation examples

Scenario: two free blocks of size 20 and 15 bytes

Allocation stream: 10, 20

- ▶ Best
- ▶ First
- ▶ Worst

Allocation stream: 8, 12, 12

- ▶ Best
- ▶ First
- ▶ Worst

Virtualizing
Memory

Memory
Virtualization

Dynamic
relocation

Segmentation

Fragmentation

28 September 2012 (18/21)

Buddy Allocation

Fast, simple allocation for blocks of 2^n bytes [Knuth]

- ▶ Raise allocation request to nearest $s = 2^n$
- ▶ Search free list for appropriate size
 - ▶ Represent free list with bitmap
 - ▶ Recursively divide larger free blocks until find free block of size s
 - ▶ “Buddy” block remains free
- ▶ Mark corresponding bits in free bitmap

To release memory:

- ▶ Mark bits as free
- ▶ Recursively coalesce freed block with buddy (if buddy is free)
 - ▶ May lazily coalesce to defer overhead

Virtualizing
Memory

Memory
Virtualization

Dynamic
relocation

Segmentation

Fragmentation

28 September 2012 (19/21)

Buddy Allocation Example

Scenario: 1 MB of free memory

Request stream:

- ▶ Allocate 70 KB, 35 KB, 80 KB
- ▶ Free 35 KB, 80 KB, 70 KB

Virtualizing
Memory

Memory
Virtualization

Dynamic
relocation

Segmentation

Fragmentation

28 September 2012 (20/21)

Comparison of Allocation Strategies

No optimal algorithm

- ▶ Fragmentation highly dependent on workload

Best fit

- ▶ Tends to leave some very large holds and some very small holes; can't use small holes easily

First fit

- ▶ Tends to leave “average” sized holes
- ▶ Advantage: faster than best fit
- ▶ Next fit used often in practice

Buddy allocation

- ▶ Minimizes external fragmentation
- ▶ Disadvantage: internal fragmentation when requests are not 2^n

Virtualizing
Memory

Memory
Virtualization

Dynamic
relocation

Segmentation

Fragmentation